

Claudio Fanconi

[✉ fanconic@gmail.com](mailto:fanconic@gmail.com) | [🏡 fanconic.github.io](https://fanconic.github.io) | [👤 fanconic](https://fanconic.net) | [📠 claudio-fanconi](https://claudio-fanconi.net) | [✉ google scholar](https://scholar.google.com/citations?user=9XzJyQIAAAJ&hl=en)

Education

University of Cambridge

PH.D. ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING

Cambridge, United Kingdom

4/2024 - 10/2027 (expected)

- Supervisor: Mihaela van der Schaar
- Researching LLM reasoning, (inverse) reinforcement learning, and alignment for decision-making
- Publications: 2 NeurIPS first-author, 1 ICML workshop first-author oral presentation
- Published paper topics: LLM model efficiency, LLM preference learning, personalised alignment

ETH Zürich

Zürich, Switzerland

M.Sc. INFORMATION TECHNOLOGY & ELECTRICAL ENGINEERING

9/2020 - 10/2022

- Focus: Machine learning and signal processing with Master thesis at **Stanford University** (3/2022 - 10/2022)
- Publications: 1 EMNLP first-author, 1 first-author journal, 1 ICML workshop first-author
- Published paper topics: Deep learning interpretability, Bayesian uncertainty estimation, NLP for healthcare
- GPA: 5.53/6.00

ETH Zürich

Zürich, Switzerland

B.Sc. INFORMATION TECHNOLOGY & ELECTRICAL ENGINEERING

9/2016 - 8/2019

- Studied abroad at the **Chinese University of Hong Kong** during the autumn semester of 2018
- Relevant subjects: datastructures & algorithms, probability & statistics, calculus, linear algebra, machine learning

Professional Experience

Sony AI

Zürich, Switzerland

RESEARCH SCIENTIST INTERN

3/2023 - 12/2023

- Collected, processed, and labelled data set (200K+ samples) to train and evaluate models for computer vision tasks
- Researched and fine-tuned deep learning models for object detection, achieving < 2px error, for robotic perception.
- Deployed proposed model in a system using TensorRT in C++ to achieve real-time predictions with sub-2ms latency

McKinsey & Company

Zürich, Switzerland

MANAGEMENT CONSULTING INTERN

6/2020 - 8/2020

- Implemented an algorithmic pricing engine for a retail company in Python to price 60K products quarterly
- Analysed, visualised, and engineered features from dataset (1M+ entries) of retail transactions using Python and SQL
- Developed impact scenarios of sales revenues, presented and explained these findings to non-technical stakeholders

IBM

Böblingen, Germany

MACHINE LEARNING ENGINEER INTERN

10/2019 - 03/2020

- Prototyped near-duplicate text detection using TensorFlow neural networks, reducing matching time by 50%.
- Implemented a data extraction tool from an MDM system in Java with HTTP requests and wrote unit tests
- Co-invented five patents on near-duplicate detection and data quality ([Google Patents](#))

Skills

Programming: Python (proficient), Java (advanced), C++ (intermediate), SQL (intermediate)

ML Frameworks: PyTorch (proficient), SciKit-Learn (proficient), TensorFlow (advanced), PyMC (advanced)

Languages: German (native), Italian (native), English (proficient), Romansh (advanced), French (basic)

Preprints

C. Fanconi, N. Astorga, M. van der Schaar - 2025

Learning Reasoning Reward Models from Expert Demonstration via Inverse Reinforcement Learning
Submitted to ICML 2026 ([Paper](#), [GitHub](#))

Conference Publications

C. Fanconi, M. van der Schaar - 2025

Cascaded Language Models for Cost-effective Human-AI Decision-Making
NeurIPS 2025 (& ICML 2025 Workshop on Multi-Agent Systems) ([Paper](#), [GitHub](#), [YouTube](#))

C. Lu*, S. Holt*, **C. Fanconi***, A.J. Chan, J. Foerster[‡], M. van der Schaar[‡], R. T. Lange[‡] - 2024

Discovering Preference Optimization Algorithms with and for Large Language Models
NeurIPS 2024 (& ICML 2024 Workshop on AutoRL) ([Paper](#), [GitHub](#), [Blog](#)) *Equal contribution

C. Fanconi*, M. Vandenhirtz*, S. Husmann, and J. Vogt - 2023

This Reads Like That: Deep Learning for Interpretable Natural Language Processing
EMNLP 2023. ([Paper](#), [GitHub](#), [YouTube](#)) *Equal contribution

C. Fanconi, M. van Buchem, and T. Hernandez-Boussard - 2023

Natural Language Processing Methods to Identify Oncology Patients at High Risk for Acute Care with Clinical Notes
AMIA 2023 Informatics Summit ([Paper](#), [GitHub](#))

Workshop Publications

K. Kobalczyk*, **C. Fanconi***, H. Sun, M. van der Schaar - 2025

Few-shot Steerable Alignment: Adapting Rewards and LLM Policies with Neural Processes
ICML 2025 Workshop on Human Feedback for AI Alignment (**oral, top 10%**) ([Paper](#), [GitHub](#)) *Equal contribution

A. Hoffmann*, **C. Fanconi***, R. Rade*, J. Kohler - 2021

This looks like that... Does it? Shortcomings of Latent Space Prototype Interpretability in Deep Networks
ICML 2021 Workshop on Explainable AI ([Paper](#), [GitHub](#), [YouTube](#)) *Equal contribution

Journal Publications

M. van Buchem, A. de Hond, **C. Fanconi**, V. Shah, M. Schüszler, I. Kant, E. Steyerberg, and T. Hernandez-Boussard - 2024

Applying Natural Language Processing to Patient Messages to Identify Depression Concerns in Cancer Patients
Journal of the American Medical Informatics Association ([Paper](#))

A. de Hond, M. van Buchem, **C. Fanconi**, M. Roy, D. Blayney, I. Kant, E. Steyerberg, and T. Hernandez-Boussard - 2024

Predicting Depression Risk in Patients With Cancer Using Multimodal Data: Algorithm Development Study
JMIR Medical Informatics ([Paper](#), [GitHub](#))

C. Fanconi, A. de Hond, D. Peterson, A. Capodici, and T. Hernandez-Boussard - 2023

A Bayesian Approach to Predictive Uncertainty in Chemotherapy Patients at Risk of Acute Care
The Lancet eBioMedicine ([Paper](#), [GitHub](#))

F. Sanmarchi, **C. Fanconi**, D. Golinelli, D. Gori, T. Hernandez-Boussard, and A. Capodici - 2023

Predict, Diagnose, and Treat Chronic Kidney Disease with Machine Learning: a Systematic Literature Review
Journal of Nephrology ([Paper](#))

Patents

L. Bremer, J. Roesner, **C. Fanconi**, M. Oberhofer, K. Steckler - 2020

Method and System for Processing Data Records
Publication Number: 20210374525 ([Google Patents](#))

T. Stuart, B. Elasioty, **C. Fanconi**, M. Grasselt, H. Babu, Y. Salliet, R. Kern, M. Oberhofer, L. Bremer, J. Roesner, J. Woods - 2020

Method for Duplicate Determination in a Graph
Publication Number: 20210374525 ([Google Patents](#))

L. Bremer, M. Oberhofer, T. Stuart, **C. Fanconi**, J. Roesner, D. Suski - 2020

Sample Pair Selection in Entity Matching Analysis
Publication Number: 20210374525 ([Google Patents](#))